RAG Testing Sample Document

1. Company Overview

Aurora Systems is a fictional technology company founded in 2018. The company develops software for document processing, artificial intelligence, workflow automation, and data analytics. Its headquarters are located in Bengaluru, India, while its research team operates from Singapore and Berlin.

Aurora Systems has approximately 420 employees. The company serves customers in 17 countries and processes more than 2.4 million documents each month.

The company's three primary products are:

- Aurora Search: an enterprise search platform for finding information across internal documents.
- Aurora Flow: a workflow automation platform used to automate repetitive business processes.
- Aurora Insight: an analytics system that generates reports from structured and unstructured data.

The company's internal slogan is "Find information. Understand context. Take action."

2. Product Information

Aurora Search was released in March 2021. It supports keyword search, semantic search, metadata filtering, and hybrid retrieval.

Aurora Flow was released in November 2022. It allows organizations to create automated workflows without writing traditional application code.

Aurora Insight was released in July 2024. It provides dashboards, scheduled reports, anomaly detection, and automated summaries.

A fictional product called Aurora Vault was announced internally in January 2026. Aurora Vault is designed to provide encrypted long-term document storage. It is not currently available to external customers.

3. Important Dates

The following dates are intentionally included for retrieval testing:

- January 12, 2018 — Aurora Systems was incorporated.
- September 4, 2019 — The first internal search prototype was completed.
- March 18, 2021 — Aurora Search launched.
- November 7, 2022 — Aurora Flow launched.

- July 22, 2024 — Aurora Insight launched.
- February 3, 2025 — The company opened its Singapore research office.
- January 16, 2026 — Aurora Vault entered internal development.
- August 11, 2026 — Aurora Systems announced its fictional 500-customer milestone.

4. Employee Information

The fictional CEO of Aurora Systems is Maya Deshmukh.

The Chief Technology Officer is Daniel Wu.

The Head of Research is Elena Petrova.

The Head of Customer Operations is Arjun Mehta.

The Finance Director is Sofia Martinez.

The company has four major engineering groups: Search Infrastructure, AI Platform, Workflow Engineering, and Data Systems.

Search Infrastructure maintains indexing and retrieval services. AI Platform develops embedding, reranking, and language-model integrations. Workflow Engineering develops Aurora Flow. Data Systems manages ingestion pipelines and analytics infrastructure.

5. Customer Support Policy

Aurora Systems provides customer support through email, chat, and a ticketing portal.

Standard support requests are normally answered within two business days. Priority support requests are targeted for response within four business hours.

Critical incidents are classified as incidents that prevent customers from accessing production services. Critical incidents are escalated immediately to the on-call engineering team.

A support ticket must contain the customer name, affected product, description of the problem, approximate time of occurrence, and relevant error messages.

6. Fictional Pricing

Aurora Search has three fictional plans.

The Starter plan costs $49 per user per month.

The Business plan costs $129 per user per month.

The Enterprise plan uses custom pricing based on document volume, number of users, storage requirements, and security requirements.

Aurora Flow costs $79 per user per month on the Standard plan and $199 per user per month on the Professional plan.

Aurora Insight is priced at $299 per workspace per month.

These prices are synthetic test data and should not be interpreted as real commercial pricing.

7. Security Rules

All production documents must be encrypted at rest and in transit.

Employees must use multi-factor authentication when accessing production systems.

Administrative accounts require hardware-based security keys.

Customer documents must not be copied into personal storage services.

Documents containing confidential customer information must have an access classification assigned before they are indexed.

The four fictional access levels are Public, Internal, Confidential, and Restricted.

Restricted documents may only be retrieved by users with explicit permission.

8. Document Retrieval Example

Consider the following fictional document:

Document ID: AUR-OPS-4721

Title: Singapore Research Office

Created: February 3, 2025

Owner: Elena Petrova

Classification: Internal

The Singapore research office focuses on multilingual retrieval, document understanding, and evaluation of AI systems. The office originally opened with 18 researchers and was expanded to 31 researchers during 2026.

The office has three research groups: Retrieval Quality, Multilingual AI, and Evaluation Systems.

Retrieval Quality investigates ranking algorithms and semantic search.

Multilingual AI studies information retrieval across languages including English, Hindi, German, Japanese, and Spanish.

Evaluation Systems develops automated tests for retrieval accuracy, grounding, and answer quality.

9. Technical Notes

A retrieval-augmented generation system generally contains several stages.

First, source documents are collected and processed.

Second, documents are divided into chunks.

Third, chunks are converted into vector representations using an embedding model.

Fourth, vectors are stored in a searchable index.

When a user submits a question, the system converts the question into a representation suitable for retrieval and searches for relevant chunks.

The retrieved context is then supplied to a language model.

Finally, the language model generates an answer based on the retrieved information.

For testing purposes, retrieval and generation should ideally be evaluated separately. A system can retrieve the wrong information, or it can retrieve the correct information and still generate an incorrect answer. Testing both stages independently makes debugging easier.
{"fallbackMarkdown":"(Google Cloud)","reference":{"matched_text":"","prefix":null,"start_idx":6706,"end_idx":6738,"safe_urls":["https://cloud.google.com/blog/products/ai-machine-learning/optimizing-rag-retrieval?hl=en","https://cloud.google.com/blog/products/ai-machine-learning/optimizing-rag-retrieval?hl=en&utm_source=chatgpt.com","https://www.teachyou.ai/blog/testing-rag-pipelines","https://www.teachyou.ai/blog/testing-rag-pipelines?utm_source=chatgpt.com"],"refs":[],"alt":"(Google Cloud)","prompt_text":null,"type":"grouped_webpages","items":[{"title":"RAG systems: Best practices to master evaluation for accurate and reliable AI. | Google Cloud Blog","url":"https://cloud.google.com/blog/products/ai-machine-learning/optimizing-rag-retrieval?hl=en&utm_source=chatgpt.com","attribution":"Google Cloud","pub_date":1734480000,"snippet":"","attribution_segments":null,"supporting_websites":[{"title":"Testing RAG Pipelines End to End: A Practical Guide · TeachYou Academy","url":"https://www.teachyou.ai/blog/testing-rag-pipelines?utm_source=chatgpt.com","pub_date":1782000000,"snippet":"","attribution":"Teach You"}],"refs":[{"turn_index":0,"ref_type":"search","ref_index":1},{"turn_index":0,"ref_type":"search","ref_index":2}],"hue":null,"attributions":null}],"error":null,"fallback_items":null,"style":null,"status":"done"},"showLoginRequiredCard":false}

10. Random Facts

The fictional Aurora Systems cafeteria serves approximately 1,250 meals each month.

The company's internal library contains 8,742 technical documents.

The average document processed by Aurora Search contains 3,800 words.

The largest document in the fictional test corpus contains 184,000 words.

The engineering team uses a fictional internal project code named "Blue Lantern."

Project Blue Lantern began on May 14, 2025.

The project originally had seven engineers and later expanded to 23 contributors.

The fictional project manager for Blue Lantern is Priya Nair.

11. Meeting Notes

Meeting: Quarterly Retrieval Review

Date: October 8, 2026

Location: Berlin Research Center

Attendees: Maya Deshmukh, Daniel Wu, Elena Petrova, Arjun Mehta, Priya Nair

Discussion topics included retrieval latency, chunk size, metadata filtering, multilingual queries, and evaluation datasets.

The team agreed to compare three retrieval configurations.

Configuration A uses keyword search only.

Configuration B uses vector search only.

Configuration C uses hybrid keyword and vector retrieval followed by reranking.

The experiment will use the same evaluation questions for all three configurations.

12. Test Questions

The following questions can be used to test the RAG system after uploading this PDF:

1. When was Aurora Systems incorporated?
2. Who is the fictional CEO?
3. Where is the company's headquarters?
4. When was Aurora Search launched?
5. How much does the fictional Business plan cost?
6. Who leads the Research department?
7. When did the Singapore office open?
8. How many researchers initially worked in Singapore?
9. What are the four access levels?
10. What is document AUR-OPS-4721 about?
11. What is Project Blue Lantern?
12. Who manages Project Blue Lantern?
13. What are the three Aurora products?
14. Which product was released in July 2024?
15. What information must a support ticket contain?
16. What happens during a critical incident?
17. What are the three retrieval configurations being compared?
18. What is the difference between retrieval failure and generation failure?
19. What languages does the Multilingual AI group study?
20. What is the fictional price of Aurora Insight?

21. Which documents require explicit permission for retrieval?
22. How many employees does Aurora Systems have?
23. What is the company's internal slogan?
24. What happened on January 16, 2026?
25. Who is the fictional Head of Customer Operations?

13. Deliberately Missing Information

The following questions should NOT have answers in this document:

- What is Aurora Systems' real revenue?
- Who is the real founder of Aurora Systems?
- What is the company's real stock price?
- What is the phone number of the Singapore office?
- What is Aurora Systems' actual website?
- How many customers does Aurora Systems have in the United States?
- What is the CEO's personal email address?

A correctly grounded RAG system should indicate that these facts are not available in the supplied document rather than inventing answers.

14. Repeated Concepts for Retrieval Testing

Retrieval testing should also include repeated concepts. The term "hybrid retrieval" appears in multiple sections because it is an important technical concept.

Hybrid retrieval combines different search approaches to improve the chance of finding relevant information. In this fictional system, hybrid retrieval means combining keyword matching with vector similarity.

Vector retrieval is useful when a question uses different wording from the source document.

Keyword retrieval can be useful when the query contains exact names, identifiers, product names, or document IDs.

Reranking can be applied after initial retrieval to reorder candidate passages according to their relevance.

These techniques should be evaluated using a fixed collection of test questions so that changes in the retrieval pipeline can be measured consistently. □{"fallbackMarkdown":"(Google Cloud)","reference":
{"matched_text":"","prefix":null,"start_idx":10702,"end_idx":10721,"safe_urls":
["https://cloud.google.com/blog/products/ai-machine-learning/optimizing-rag-retrieval?hl=en","https://cloud.google.com/blog/products/ai-machine-learning/optimizing-rag-retrieval?hl=en&utm_source=chatgpt.com"],"refs":[],"alt":"(Google

Cloud)","prompt_text":null,"type":"grouped_webpages","items":[{"title":"RAG systems: Best practices to master evaluation for accurate and reliable AI. | Google Cloud Blog","url":"https://cloud.google.com/blog/products/ai-machine-learning/optimizing-rag-retrieval?hl=en&utm_source=chatgpt.com","attribution":"Google Cloud","pub_date":1734480000,"snippet":"","attribution_segments":null,"supporting_websites":[],"refs":[{"turn_index":0,"ref_type":"search","ref_index":1}],"hue":null,"attributions":null}],"error":null,"fallback_items":null,"style":null,"status":"done"},"showLoginRequiredCard":false}